

I Flare CPU Pipeline Stalls With PipeMemRmw

I.A Pipeline Stages EX and MEM

I.A.1 Loads/Stores/Atomics with Memory

- MEM should stall when there's a load/store/atomic-with-memory instruction that has a dcache miss. This means that EX needs to stall too.
- Here are the main situations of this for when forwarding loaded values (including atomic swap-type instructions) results from MEM to EX.
 - `xchg.l/cmpxchg/cmpxchg.l` (which are my only non-load/store/stack access instructions with memory arguments)
 - The atomic forms of these instructions, i.e. `xchg.l` and `cmpxchg.l`, will need to stall MEM while the bus is accessed. Treat these similarly to **load/store** instructions that have dcache misses, i.e. stall until we can forward.
 - These instructions each do one load and one store
 - `push/pop`:
 - `pop` at least needs to store new values into two registers. How do I handle multiple register updates? `PipeMemRmw` only supports one backing-mem write per transaction. The best I can figure is to stall ID upon such an instruction being decoded until it finishes. That is what I guess I'll do. Unfortunately this makes `pop` take more cycles even upon a dcache hit, but it's the only thing I can think of for how to handle memory-access instructions that drive more than one register. At least `push`, for register writes, only affects the register used as a stack pointer, which will *usually* be `sp`, but that is not enforced by the instruction set at all. If I were making an OoO version of the Flare CPU, maybe it'd be wise to optimize the case of using `sp` for the `push/pop` stack pointer operand somehow.
 - It may make sense to just get rid of `push/pop` as instructions since a combination of `ldr/str` and an `add sp, #...` instruction would possibly be faster if the addresses that are accessed are in dcache
- EX has a **load**-based RAW pipeline hazard with an instruction that doesn't stall outside of this situation.
 - Example code:

```
ldr r1, [r2]
add r3, r1
```
 - Stalling states:
 - EX: stalls until it can forward from MEM.
 - MEM: stalls until it finishes the **load** instruction.
 - EX checks for a **load** instruction and changes FSM state upon `EX.up.isFiring == True`.
 - What happens if after `EX.up.isFiring` that there's an instance of `EX.haltIt()`? Does MEM still see the instruction that was previously in EX? The answer should be **yes**.
 - If there's a dcache miss, MEM should stall by doing `MEM.duplicateIt()` and also deassert `MEM.up(...).modMemWordValid`
 - Also, stop `MEM.up(...).modMemWordValid` upon no longer doing `MEM.duplicateIt()`.
 - If there's a dcache hit, MEM should not stall and I think that EX doesn't need to either? That's cool! I thought there'd be a hazard even in this case, but if the address loaded from was in dcache, I can already forward, as I already have the result in that case (i.e. I already read it from the BRAM).
- EX has a **load**-based RAW pipeline hazard with an instruction that **does** stall normally in EX, i.e. `mul`, `div`, etc.
 - I can't have `MEM.up.isFiring` if I do `EX.haltIt()` immediately following a detected **load** instruction. Thus I think I should do `EX.duplicateIt()` but with pipeline bubbles sent from EX to MEM until `MEM.up.isFiring/EX.down.isFiring`.

- Pipeline bubbles should be something other than `modMemWordValid` in the actual Flare CPU pipeline, as `modMemWordValid` is only for `PipeMemRmw` itself.
- Upon `EX.down.isFiring` (and following the detected **load** in `EX.up.isFiring`), `EX` should save a copy of the data being forwarded from `MEM`, then on the following cycle `EX` can send it to the multiplier, divider, etc. For those following cycles until `EX.up.isFiring`, `EX` should do `EX.haltIt()` until the mul, div, etc. completes. On the cycle following the completed instruction, don't do `EX.haltIt()` and maintain the value of `tempUpMod(2)` until `EX.up.isFiring`.
 - I have existing code in the `PipeMemRmw` tester stuff that already does this kind of capture, in the fake `MEM` stage, but `EX` in the Flare CPU needs to do the same thing with its own instruction.
- `PipeMemRmw` already checks for `up.isValid` when capturing the input pipe payload into `upExt(1)(extIdxSingle)`, so we can at least see the instruction via `upExt(1)(extIdxSingle)` in the `doModInModFrontFunc()` generic. Actually, `doModInModFront()` takes `tempUpMod(1)`, which does still contain the value of `upExt(1)(extIdxSingle)`.
- When `EX` stalls, *maybe* it should not change the value of `tempUpMod(2)`? Instead *maybe* `EX` should change the value of `tempUpMod(2)` upon no longer stalling via `EX.duplicateIt()`. `tempUpMod(2)` should have its value set to the output of `EX` upon the multiply/divide finishing, at which point we stop doing `EX.duplicateIt()`, after which `tempUpMod(2)`'s value should be held constant until `EX.up.isFiring`.
- Is everything similar for **store** instructions but without the forwarding? I guess just do everything that was needed for a **load** instruction besides the forwarding. Thus deassert a **store** instruction's `modMemWordValid` to prevent its "results" from being forwarded. This approach is guaranteed to work and will use less logic than specially handling **stores**.
- **Is there actually no hazard here? Other than that we can't forward upon a dcache miss?**

I.B